

Data Science: Capstone - Create Your Own Project (Prediction of Breast Cancer)

KLC

May 1, 2020

Objective

The aim of this project is to build a machine learning algorithm to predict whether a breast mass cell is “benign” or “malignant”. Breast lumps that are benign are mostly non-cancerous and not life threatening. They do not spread outside of the breast. Malignant lumps are cancerous however. Such kind of prediction algorithm could help medical practitioners to detect and diagnose breast cancer.

We will train a few machine learning models and measure their performance with their prediction sensitivity and F1 score. Our aim is to choose a model which yields the highest sensitivity (i.e. low false-negative) and F1 score.

This report will first explore the dataset, then analyse several models, compare their performance and conclude the result.

Dataset

Breast Cancer Wisconsin (Diagnostic) DataSet obtained from Kaggle (<https://www.kaggle.com/lbronchal/breast-cancer-dataset-analysis>) is used for this project. Such data, collected in 1993 by the University of Wisconsin, contains 569 samples of measurements on cells in suspicious lumps in a women’s breast. 20% of the data will be used for testing, while the remaining will be used for training the machine learning algorithm.

Description of data: The dataset contains 569 observations with 33 variables including 30 “features” as listed below. “Features” are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass, which describe characteristics of the cell nuclei present in the image.

Attribute Information:

- 1) ID number
- 2) Diagnosis (M = malignant, B = benign) 3-32)

Ten real-valued features are computed for each cell nucleus:

- a) radius (mean of distances from center to points on the perimeter)
- b) texture (standard deviation of gray-scale values)
- c) perimeter
- d) area
- e) smoothness (local variation in radius lengths)
- f) compactness ($\text{perimeter}^2 / \text{area} - 1.0$)
- g) concavity (severity of concave portions of the contour)
- h) concave points (number of concave portions of the contour)
- i) symmetry
- j) fractal dimension (“coastline approximation” - 1)

The mean, standard error and “worst” or largest (mean of the three largest values) of these features were computed for each image, resulting in 30 features. All feature values are recorded with 4 significant digits.

```
#####  
# Load dataset  
#####  
  
data <- read.csv("https://raw.githubusercontent.com/happycheers/BreastCancer/master/data.csv")
```

Data exploration

First, let's grab an overview of the dataset.

```
#Structure of the dataset  
str(data)
```

```
## 'data.frame':   569 obs. of  33 variables:  
## $ id           : int  842302 842517 84300903 84348301 84358402 843786 844359 84458202 844...  
## $ diagnosis    : Factor w/ 2 levels "B","M": 2 2 2 2 2 2 2 2 2 ...  
## $ radius_mean  : num  18 20.6 19.7 11.4 20.3 ...  
## $ texture_mean : num  10.4 17.8 21.2 20.4 14.3 ...  
## $ perimeter_mean : num  122.8 132.9 130 77.6 135.1 ...  
## $ area_mean    : num  1001 1326 1203 386 1297 ...  
## $ smoothness_mean : num  0.1184 0.0847 0.1096 0.1425 0.1003 ...  
## $ compactness_mean : num  0.2776 0.0786 0.1599 0.2839 0.1328 ...  
## $ concavity_mean : num  0.3001 0.0869 0.1974 0.2414 0.198 ...  
## $ concave.points_mean : num  0.1471 0.0702 0.1279 0.1052 0.1043 ...  
## $ symmetry_mean  : num  0.242 0.181 0.207 0.26 0.181 ...  
## $ fractal_dimension_mean : num  0.0787 0.0567 0.06 0.0974 0.0588 ...  
## $ radius_se      : num  1.095 0.543 0.746 0.496 0.757 ...  
## $ texture_se      : num  0.905 0.734 0.787 1.156 0.781 ...  
## $ perimeter_se    : num  8.59 3.4 4.58 3.44 5.44 ...  
## $ area_se         : num  153.4 74.1 94 27.2 94.4 ...  
## $ smoothness_se   : num  0.0064 0.00522 0.00615 0.00911 0.01149 ...  
## $ compactness_se  : num  0.049 0.0131 0.0401 0.0746 0.0246 ...  
## $ concavity_se    : num  0.0537 0.0186 0.0383 0.0566 0.0569 ...  
## $ concave.points_se : num  0.0159 0.0134 0.0206 0.0187 0.0188 ...  
## $ symmetry_se      : num  0.03 0.0139 0.0225 0.0596 0.0176 ...  
## $ fractal_dimension_se : num  0.00619 0.00353 0.00457 0.00921 0.00511 ...  
## $ radius_worst     : num  25.4 25 23.6 14.9 22.5 ...  
## $ texture_worst    : num  17.3 23.4 25.5 26.5 16.7 ...  
## $ perimeter_worst  : num  184.6 158.8 152.5 98.9 152.2 ...  
## $ area_worst       : num  2019 1956 1709 568 1575 ...  
## $ smoothness_worst : num  0.162 0.124 0.144 0.21 0.137 ...  
## $ compactness_worst : num  0.666 0.187 0.424 0.866 0.205 ...  
## $ concavity_worst  : num  0.712 0.242 0.45 0.687 0.4 ...  
## $ concave.points_worst : num  0.265 0.186 0.243 0.258 0.163 ...  
## $ symmetry_worst   : num  0.46 0.275 0.361 0.664 0.236 ...  
## $ fractal_dimension_worst : num  0.1189 0.089 0.0876 0.173 0.0768 ...  
## $ X                : logi  NA NA NA NA NA NA ...
```

```
# First 6 rows and header  
head(data)
```

```
##           id diagnosis radius_mean texture_mean perimeter_mean area_mean
```

```

## 1 842302 M 17.99 10.38 122.80 1001.0
## 2 842517 M 20.57 17.77 132.90 1326.0
## 3 84300903 M 19.69 21.25 130.00 1203.0
## 4 84348301 M 11.42 20.38 77.58 386.1
## 5 84358402 M 20.29 14.34 135.10 1297.0
## 6 843786 M 12.45 15.70 82.57 477.1
## smoothness_mean compactness_mean concavity_mean concave.points_mean
## 1 0.11840 0.27760 0.3001 0.14710
## 2 0.08474 0.07864 0.0869 0.07017
## 3 0.10960 0.15990 0.1974 0.12790
## 4 0.14250 0.28390 0.2414 0.10520
## 5 0.10030 0.13280 0.1980 0.10430
## 6 0.12780 0.17000 0.1578 0.08089
## symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 1 0.2419 0.07871 1.0950 0.9053 8.589
## 2 0.1812 0.05667 0.5435 0.7339 3.398
## 3 0.2069 0.05999 0.7456 0.7869 4.585
## 4 0.2597 0.09744 0.4956 1.1560 3.445
## 5 0.1809 0.05883 0.7572 0.7813 5.438
## 6 0.2087 0.07613 0.3345 0.8902 2.217
## area_se smoothness_se compactness_se concavity_se concave.points_se
## 1 153.40 0.006399 0.04904 0.05373 0.01587
## 2 74.08 0.005225 0.01308 0.01860 0.01340
## 3 94.03 0.006150 0.04006 0.03832 0.02058
## 4 27.23 0.009110 0.07458 0.05661 0.01867
## 5 94.44 0.011490 0.02461 0.05688 0.01885
## 6 27.19 0.007510 0.03345 0.03672 0.01137
## symmetry_se fractal_dimension_se radius_worst texture_worst perimeter_worst
## 1 0.03003 0.006193 25.38 17.33 184.60
## 2 0.01389 0.003532 24.99 23.41 158.80
## 3 0.02250 0.004571 23.57 25.53 152.50
## 4 0.05963 0.009208 14.91 26.50 98.87
## 5 0.01756 0.005115 22.54 16.67 152.20
## 6 0.02165 0.005082 15.47 23.75 103.40
## area_worst smoothness_worst compactness_worst concavity_worst
## 1 2019.0 0.1622 0.6656 0.7119
## 2 1956.0 0.1238 0.1866 0.2416
## 3 1709.0 0.1444 0.4245 0.4504
## 4 567.7 0.2098 0.8663 0.6869
## 5 1575.0 0.1374 0.2050 0.4000
## 6 741.6 0.1791 0.5249 0.5355
## concave.points_worst symmetry_worst fractal_dimension_worst X
## 1 0.2654 0.4601 0.11890 NA
## 2 0.1860 0.2750 0.08902 NA
## 3 0.2430 0.3613 0.08758 NA
## 4 0.2575 0.6638 0.17300 NA
## 5 0.1625 0.2364 0.07678 NA
## 6 0.1741 0.3985 0.12440 NA

```

Summary of statistics

```
summary(data)
```

```

##      id      diagnosis radius_mean texture_mean
## Min.   :    8670 B:357   Min.   : 6.981   Min.   : 9.71
## 1st Qu.:   869218 M:212   1st Qu.:11.700 1st Qu.:16.17

```

```

## Median : 906024          Median :13.370    Median :18.84
## Mean : 30371831         Mean :14.127    Mean :19.29
## 3rd Qu.: 8813129        3rd Qu.:15.780    3rd Qu.:21.80
## Max. :911320502         Max. :28.110    Max. :39.28
## perimeter_mean    area_mean    smoothness_mean    compactness_mean
## Min. : 43.79      Min. : 143.5      Min. :0.05263      Min. :0.01938
## 1st Qu.: 75.17    1st Qu.: 420.3    1st Qu.:0.08637    1st Qu.:0.06492
## Median : 86.24    Median : 551.1    Median :0.09587    Median :0.09263
## Mean : 91.97      Mean : 654.9      Mean :0.09636      Mean :0.10434
## 3rd Qu.:104.10    3rd Qu.: 782.7    3rd Qu.:0.10530    3rd Qu.:0.13040
## Max. :188.50      Max. :2501.0      Max. :0.16340      Max. :0.34540
## concavity_mean    concave.points_mean symmetry_mean    fractal_dimension_mean
## Min. :0.00000      Min. :0.00000      Min. :0.1060      Min. :0.04996
## 1st Qu.:0.02956    1st Qu.:0.02031    1st Qu.:0.1619    1st Qu.:0.05770
## Median :0.06154    Median :0.03350    Median :0.1792    Median :0.06154
## Mean :0.08880      Mean :0.04892      Mean :0.1812      Mean :0.06280
## 3rd Qu.:0.13070    3rd Qu.:0.07400    3rd Qu.:0.1957    3rd Qu.:0.06612
## Max. :0.42680      Max. :0.20120      Max. :0.3040      Max. :0.09744
## radius_se         texture_se         perimeter_se         area_se
## Min. :0.1115      Min. :0.3602      Min. : 0.757      Min. : 6.802
## 1st Qu.:0.2324    1st Qu.:0.8339    1st Qu.: 1.606    1st Qu.: 17.850
## Median :0.3242    Median :1.1080    Median : 2.287    Median : 24.530
## Mean :0.4052      Mean :1.2169      Mean : 2.866      Mean : 40.337
## 3rd Qu.:0.4789    3rd Qu.:1.4740    3rd Qu.: 3.357    3rd Qu.: 45.190
## Max. :2.8730      Max. :4.8850      Max. :21.980      Max. :542.200
## smoothness_se      compactness_se      concavity_se      concave.points_se
## Min. :0.001713      Min. :0.002252      Min. :0.00000      Min. :0.000000
## 1st Qu.:0.005169    1st Qu.:0.013080    1st Qu.:0.01509    1st Qu.:0.007638
## Median :0.006380    Median :0.020450    Median :0.02589    Median :0.010930
## Mean :0.007041      Mean :0.025478      Mean :0.03189      Mean :0.011796
## 3rd Qu.:0.008146    3rd Qu.:0.032450    3rd Qu.:0.04205    3rd Qu.:0.014710
## Max. :0.031130      Max. :0.135400      Max. :0.39600      Max. :0.052790
## symmetry_se         fractal_dimension_se radius_worst      texture_worst
## Min. :0.007882      Min. :0.0008948      Min. : 7.93      Min. :12.02
## 1st Qu.:0.015160    1st Qu.:0.0022480    1st Qu.:13.01    1st Qu.:21.08
## Median :0.018730    Median :0.0031870    Median :14.97    Median :25.41
## Mean :0.020542      Mean :0.0037949      Mean :16.27      Mean :25.68
## 3rd Qu.:0.023480    3rd Qu.:0.0045580    3rd Qu.:18.79    3rd Qu.:29.72
## Max. :0.078950      Max. :0.0298400      Max. :36.04      Max. :49.54
## perimeter_worst     area_worst          smoothness_worst    compactness_worst
## Min. : 50.41      Min. : 185.2      Min. :0.07117      Min. :0.02729
## 1st Qu.: 84.11    1st Qu.: 515.3    1st Qu.:0.11660    1st Qu.:0.14720
## Median : 97.66    Median : 686.5    Median :0.13130    Median :0.21190
## Mean :107.26      Mean : 880.6      Mean :0.13237      Mean :0.25427
## 3rd Qu.:125.40    3rd Qu.:1084.0    3rd Qu.:0.14600    3rd Qu.:0.33910
## Max. :251.20      Max. :4254.0      Max. :0.22260      Max. :1.05800
## concavity_worst     concave.points_worst symmetry_worst      fractal_dimension_worst
## Min. :0.0000      Min. :0.00000      Min. :0.1565      Min. :0.05504
## 1st Qu.:0.1145    1st Qu.:0.06493    1st Qu.:0.2504    1st Qu.:0.07146
## Median :0.2267      Median :0.09993      Median :0.2822      Median :0.08004
## Mean :0.2722      Mean :0.11461      Mean :0.2901      Mean :0.08395
## 3rd Qu.:0.3829      3rd Qu.:0.16140      3rd Qu.:0.3179      3rd Qu.:0.09208
## Max. :1.2520      Max. :0.29100      Max. :0.6638      Max. :0.20750
## X

```

```
## Mode:logical
## NA's:569
##
##
##
##

# Summarize number of diagnosis ("B" and "M") in the dataset
data %>% group_by(diagnosis) %>% summarize(n())

## # A tibble: 2 x 2
##   diagnosis `n()`
##   <fct>      <int>
## 1 B          357
## 2 M          212
```

We note from the above that the diagnosis is slightly unbalanced. We may have to adjust the data when using some models so that they can work properly.

Data cleaning

```
#####
# Data cleaning
#####

# Remove columns 1 and 33 as irrelevant
data <- data[,-33]
data <- data[,-1]

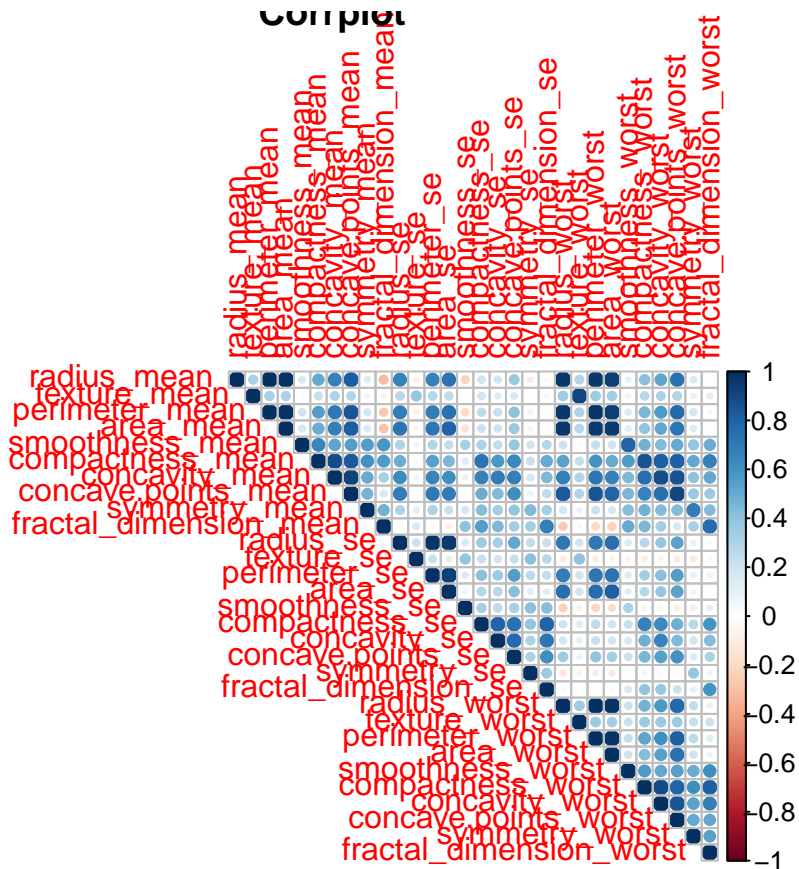
# Check if there are missing values
map_int(data, function(.x) sum(is.na(.x)))

##           diagnosis      radius_mean      texture_mean
##           0              0              0
##   perimeter_mean      area_mean      smoothness_mean
##           0              0              0
##   compactness_mean      concavity_mean      concave.points_mean
##           0              0              0
##   symmetry_mean      fractal_dimension_mean      radius_se
##           0              0              0
##   texture_se      perimeter_se      area_se
##           0              0              0
##   smoothness_se      compactness_se      concavity_se
##           0              0              0
##   concave.points_se      symmetry_se      fractal_dimension_se
##           0              0              0
##   radius_worst      texture_worst      perimeter_worst
##           0              0              0
##   area_worst      smoothness_worst      compactness_worst
##           0              0              0
##   concavity_worst      concave.points_worst      symmetry_worst
##           0              0              0
##   fractal_dimension_worst
##           0
```

We now have removed the 1st column “id” and the 33rd column “X” as they appear irrelevant to our prediction. We have also checked that there are no missing values in the dataset.

Some models such as naive bayes do not work well with highly-correlated variables as they assume the predictor variables are independent with each other. Therefore, we will check if the variables of the dataset are highly correlated and will remove them if their correlation coefficients are higher than 0.9 or lower than -0.9.

```
# Plot the correlation among variables
corrplot(cor(data[,2:31]) , main=" Corrplot" , method = "circle" , type = "upper")
```



```
# Identify variables with correlation coefficient higher than 0.9 or lower than -0.9
to_drop_col <- findCorrelation(cor(data[,2:31]), cutoff=0.9)
```

```
# Adjust the result by one column shift
to_drop_col <- to_drop_col + 1
```

```
# Remove highly correlated variables
new_data <- data[, -to_drop_col]
```

```
# Cross-check if highly correlated variables have been removed
findCorrelation(cor(new_data[,2:21]), cutoff=0.9)
```

```
## integer(0)
```

Now, we are going to divide the dataset into training (80%) and testing (20%) datasets.

```
#####
# Create training and testing sets
```

```
#####

# Divide the data set into training (80%) and testing (20%) sets
set.seed(1234, sample.kind="Rounding")

## Warning in set.seed(1234, sample.kind = "Rounding"): non-uniform 'Rounding'
## sampler used

index <- createDataPartition(new_data$diagnosis, times=1, p=0.8, list = FALSE)
train <- new_data[index, ]
test <- new_data[-index, ]
```

Data Analysis - Modelling Approach

In the following, we will train a naive bayes model, logistic regression model, k-nearest neighbor model and random forest model.

```
# Cross validation with 10 folds
tc <- trainControl(method="cv", number = 10, classProbs=TRUE, summaryFunction = twoClassSummary)
```

Naive Bayes Model

```
#####
# Naive bayes model
#####

# Train a naive bayes model
naiveb_model <- train(diagnosis~.,
                      train,
                      method="nb",
                      metric = "ROC",
                      preProcess=c('center','scale'),
                      trControl=tc)

# Predict testing set
naiveb_pred <- predict(naiveb_model, test)

# summarize results (set positive as "M" so that the sensitivity is correct)
naiveb_result <- confusionMatrix(naiveb_pred, test$diagnosis, positive = "M")
naiveb_result
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  B  M
##           B 67  3
##           M  4 39
##
##           Accuracy : 0.9381
##           95% CI : (0.8765, 0.9747)
##           No Information Rate : 0.6283
##           P-Value [Acc > NIR] : 1.718e-14
##
##           Kappa : 0.868
```

```
##
## McNemar's Test P-Value : 1
##
##          Sensitivity : 0.9286
##          Specificity : 0.9437
##          Pos Pred Value : 0.9070
##          Neg Pred Value : 0.9571
##          Prevalence : 0.3717
##          Detection Rate : 0.3451
##          Detection Prevalence : 0.3805
##          Balanced Accuracy : 0.9361
##
##          'Positive' Class : M
##
```

Logistic Regression Model

```
#####
# Logistic regression model
#####

# Train a logistic regression model
glm_model <- train(diagnosis~.,
                  train,
                  method="glm",
                  metric = "ROC",
                  preProcess=c('center','scale'),
                  trControl=tc)

# Predict testing set
glm_pred <- predict(glm_model, test)

# summarize results (set positive as "M" so that the sensitivity is correct)
glm_result <- confusionMatrix(glm_pred, test$diagnosis, positive = "M")
glm_result
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  B  M
##          B 69  1
##          M  2 41
##
##          Accuracy : 0.9735
##          95% CI : (0.9244, 0.9945)
##          No Information Rate : 0.6283
##          P-Value [Acc > NIR] : <2e-16
##
##          Kappa : 0.9434
##
## McNemar's Test P-Value : 1
##
##          Sensitivity : 0.9762
##          Specificity : 0.9718
```



```
##          Pos Pred Value : 0.9535
##          Neg Pred Value : 0.9857
##          Prevalence : 0.3717
##          Detection Rate : 0.3628
##          Detection Prevalence : 0.3805
##          Balanced Accuracy : 0.9740
##
##          'Positive' Class : M
##
```

K-nearest Neighbor Model

```
#####
# K-nearest neighbor model
#####

# Train a KNN model
knn_model <- train(diagnosis~.,
                  train,
                  method="knn",
                  metric = "ROC",
                  preProcess=c('center','scale'),
                  tuneLength=10,
                  trControl=tc)

# Predict testing set
knn_pred <- predict(knn_model, test)

# summarize results (set positive as "M" so that the sensitivity is correct)
knn_result <- confusionMatrix(knn_pred, test$diagnosis, positive = "M")
knn_result
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  B  M
##          B 71  3
##          M  0 39
##
##          Accuracy : 0.9735
##          95% CI : (0.9244, 0.9945)
##          No Information Rate : 0.6283
##          P-Value [Acc > NIR] : <2e-16
##
##          Kappa : 0.9423
##
##          Mcnemar's Test P-Value : 0.2482
##
##          Sensitivity : 0.9286
##          Specificity : 1.0000
##          Pos Pred Value : 1.0000
##          Neg Pred Value : 0.9595
##          Prevalence : 0.3717
##          Detection Rate : 0.3451
```

```
##      Detection Prevalence : 0.3451
##      Balanced Accuracy : 0.9643
##
##      'Positive' Class : M
##
```

Random Forest Model

```
#####
# Random forest model
#####

# Train a random forest model
rf_model <- train(diagnosis~.,
                  train,
                  method="rf",
                  metric = "ROC",
                  preProcess=c('center','scale'),
                  trControl=tc)

# Predict testing set
rf_pred <- predict(rf_model, test)

# summarize results (set positive as "M" so that the sensitivity is correct)
rf_result <- confusionMatrix(rf_pred, test$diagnosis, positive = "M")
rf_result
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction B  M
##      B 67  3
##      M  4 39
##
##      Accuracy : 0.9381
##      95% CI : (0.8765, 0.9747)
##      No Information Rate : 0.6283
##      P-Value [Acc > NIR] : 1.718e-14
##
##      Kappa : 0.868
##
##      McNemar's Test P-Value : 1
##
##      Sensitivity : 0.9286
##      Specificity : 0.9437
##      Pos Pred Value : 0.9070
##      Neg Pred Value : 0.9571
##      Prevalence : 0.3717
##      Detection Rate : 0.3451
##      Detection Prevalence : 0.3805
##      Balanced Accuracy : 0.9361
##
##      'Positive' Class : M
##
```

Results

The results of each model developed above are summarized below:

```
#####  
# Results  
#####  
  
# Summarize the confusion matrixes of each model  
result_list <- list (naive_bayes = naiveb_result,  
                     logistic_regression = glm_result,  
                     KNN = knn_result,  
                     random_forest = rf_result)  
results <- sapply (result_list, function(x) x$byClass)  
  
# Print the results in a table format  
results %>% knitr::kable()
```

	naive_bayes	logistic_regression	KNN	random_forest
Sensitivity	0.9285714	0.9761905	0.9285714	0.9285714
Specificity	0.9436620	0.9718310	1.0000000	0.9436620
Pos Pred Value	0.9069767	0.9534884	1.0000000	0.9069767
Neg Pred Value	0.9571429	0.9857143	0.9594595	0.9571429
Precision	0.9069767	0.9534884	1.0000000	0.9069767
Recall	0.9285714	0.9761905	0.9285714	0.9285714
F1	0.9176471	0.9647059	0.9629630	0.9176471
Prevalence	0.3716814	0.3716814	0.3716814	0.3716814
Detection Rate	0.3451327	0.3628319	0.3451327	0.3451327
Detection Prevalence	0.3805310	0.3805310	0.3451327	0.3805310
Balanced Accuracy	0.9361167	0.9740107	0.9642857	0.9361167

```
# Identify the best results for each metric in confusion matrix  
best_results <- apply(results, 1, which.is.max)  
  
# Match the best results with corresponding model  
report <- tibble (metric = names(best_results),  
                  best_model = colnames(results)[best_results],  
                  value=mapply(function(x,y) {results[x,y]},  
                                names(best_results),  
                                best_results))  
rownames(report)<-NULL  
  
# Print the best model identified for each metric  
report
```

```
## # A tibble: 11 x 3  
##   metric          best_model      value  
##   <chr>          <chr>          <dbl>  
## 1 Sensitivity    logistic_regression 0.976  
## 2 Specificity    KNN              1  
## 3 Pos Pred Value KNN              1  
## 4 Neg Pred Value logistic_regression 0.986  
## 5 Precision      KNN              1  
## 6 Recall         logistic_regression 0.976
```

## 7 F1	logistic_regression	0.965
## 8 Prevalence	KNN	0.372
## 9 Detection Rate	logistic_regression	0.363
## 10 Detection Prevalence	random_forest	0.381
## 11 Balanced Accuracy	logistic_regression	0.974

There has been discussion which metric, say accuracy, precision, recall or F1 score, we should use to select the “best model”. There is no one-size-fit-all answer. **For our case on predicting whether a breast mass cell is cancerous, undoubtedly the cost associated with false negative is high.** The consequence can be very serious for a patient where his/her cell is predicted as negative (benign) while it is actually positive (malignant). In this regard, **the sensitivity, which calculates how many of actual positives a model capture, would be more relevant in this case.** While false positives appear to cause less serious consequence in our case, it does not mean costless. A healthy person diagnosed with cancer will result in stress and high medical costs. To this end, we should consider F1 score as well, which seek a balance between sensitivity and specificity.

Based on the result table above, the **best model should be logistic regression model** which has the highest sensitivity and F1 score.

Conclusion

In this project, We have developed four machine learning models to predict classification of a breast mass cell as “benign” or “malignant”. Then, we have discussed which metric we should use to select the best model. Finally, we have selected logistic regression model as the most optimal one given its good sensitivity and F1 score. To further improve our prediction, we can in fact build more models such as neural network and support vector machine to explore if there are any better models than the one we chose.

This kind of classification prediction will have a wide use across industries, such as predicting no-shows for medical appointments, spam emails or fraudulent transactions.